

Machine Learning & Spark

MACHINE LEARNING WITH PYSPARK



Andrew Collier

Data Scientist, Fathom Data

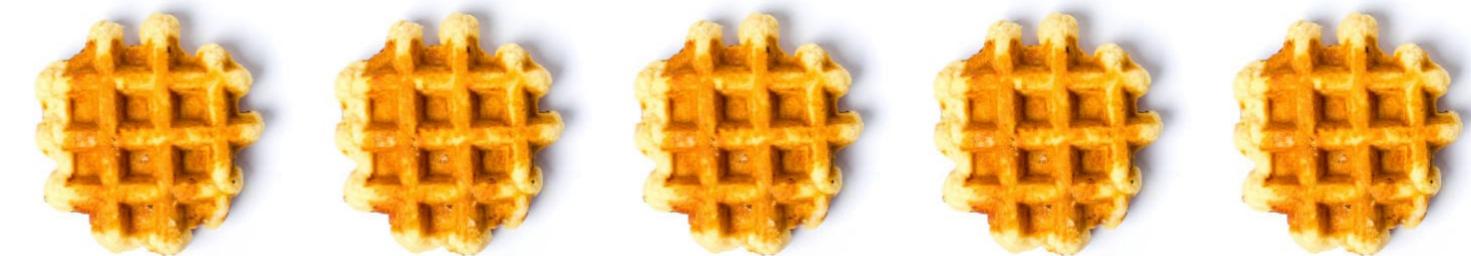
Building the perfect waffle (an analogy)



Archetype Waffle

Find waffle recipe. Give explicit instructions:

- 125 g flour
- 1 t baking powder
- 1 egg
- 225 ml milk
- 1 T melted butter



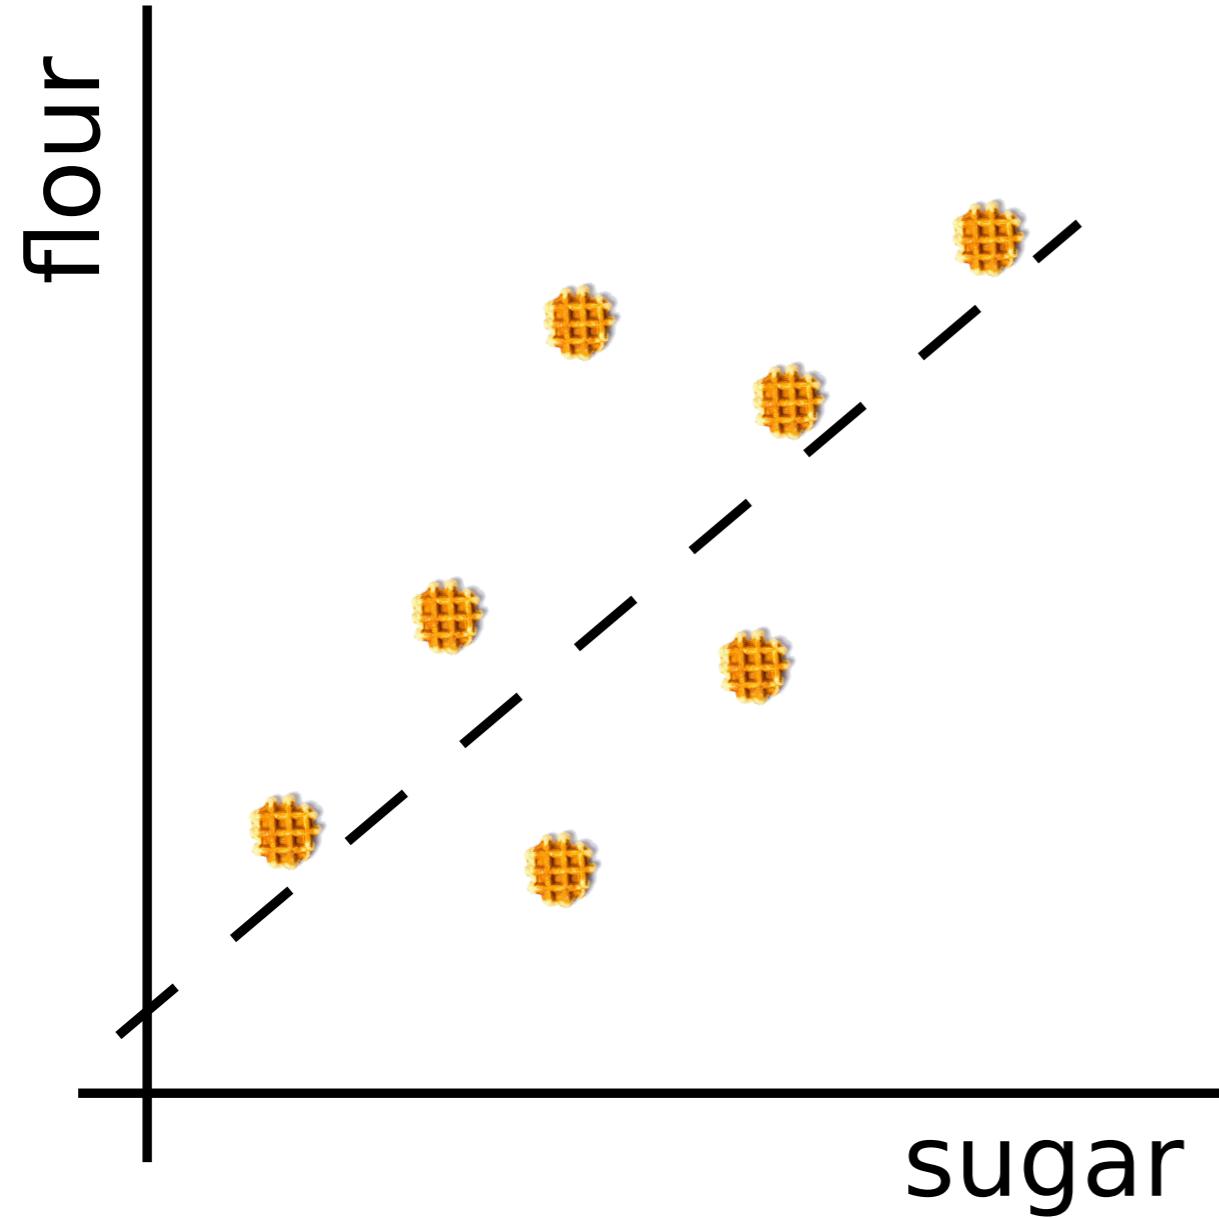
Find many waffle recipes.

Learn the perfect recipe:

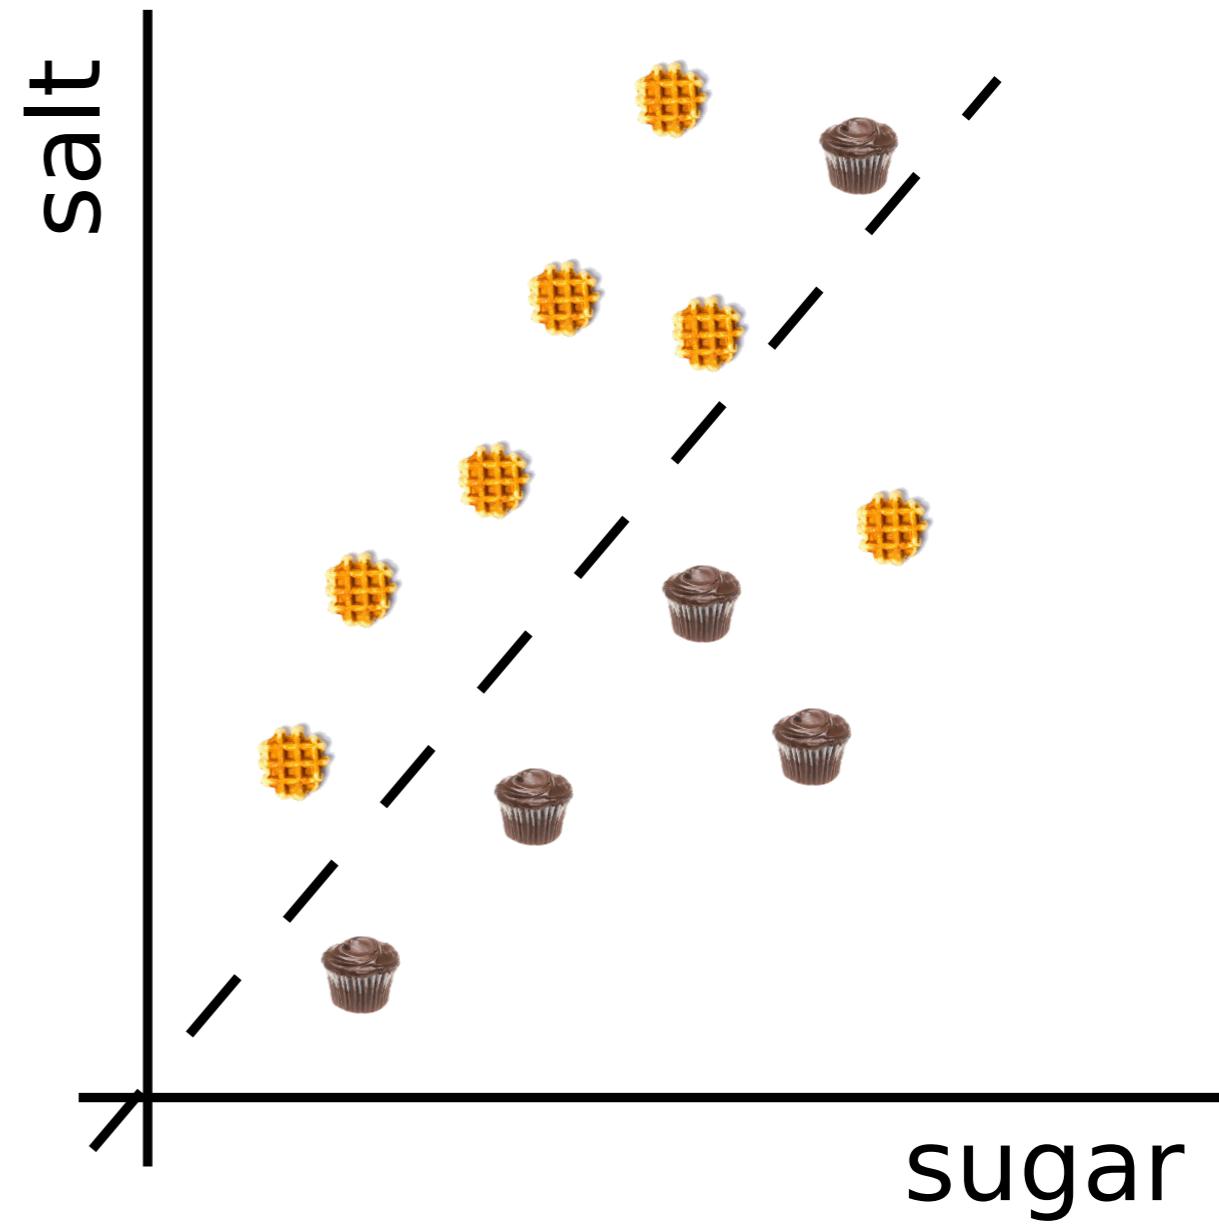
1. Look at lots of recipes.
2. What ingredients?
3. What proportions?

Computer generates its own instructions.

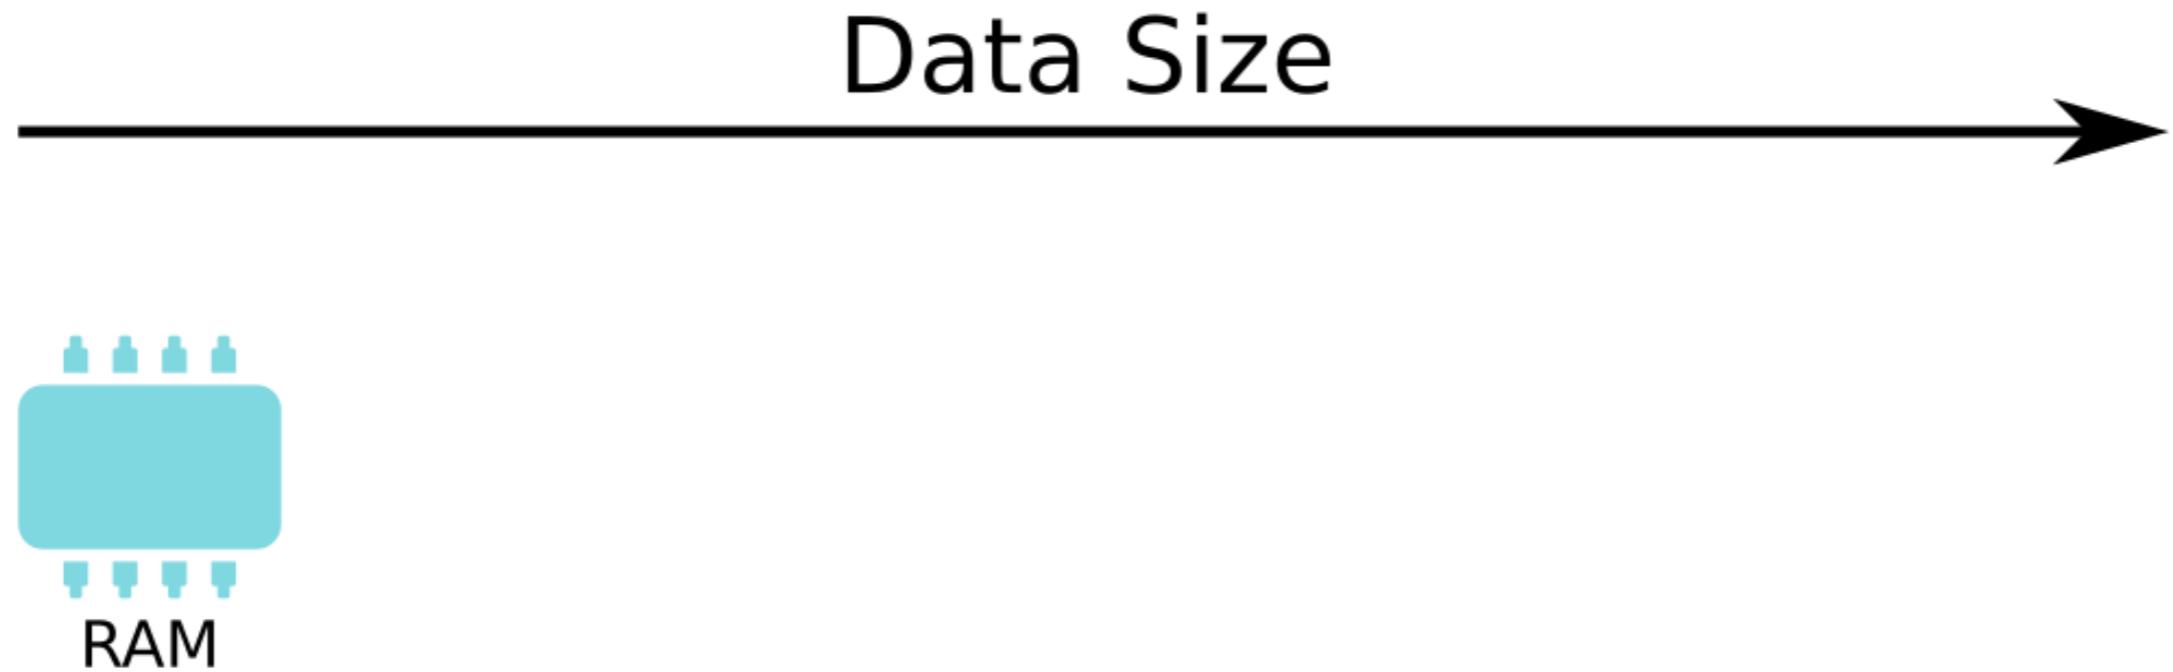
Regression



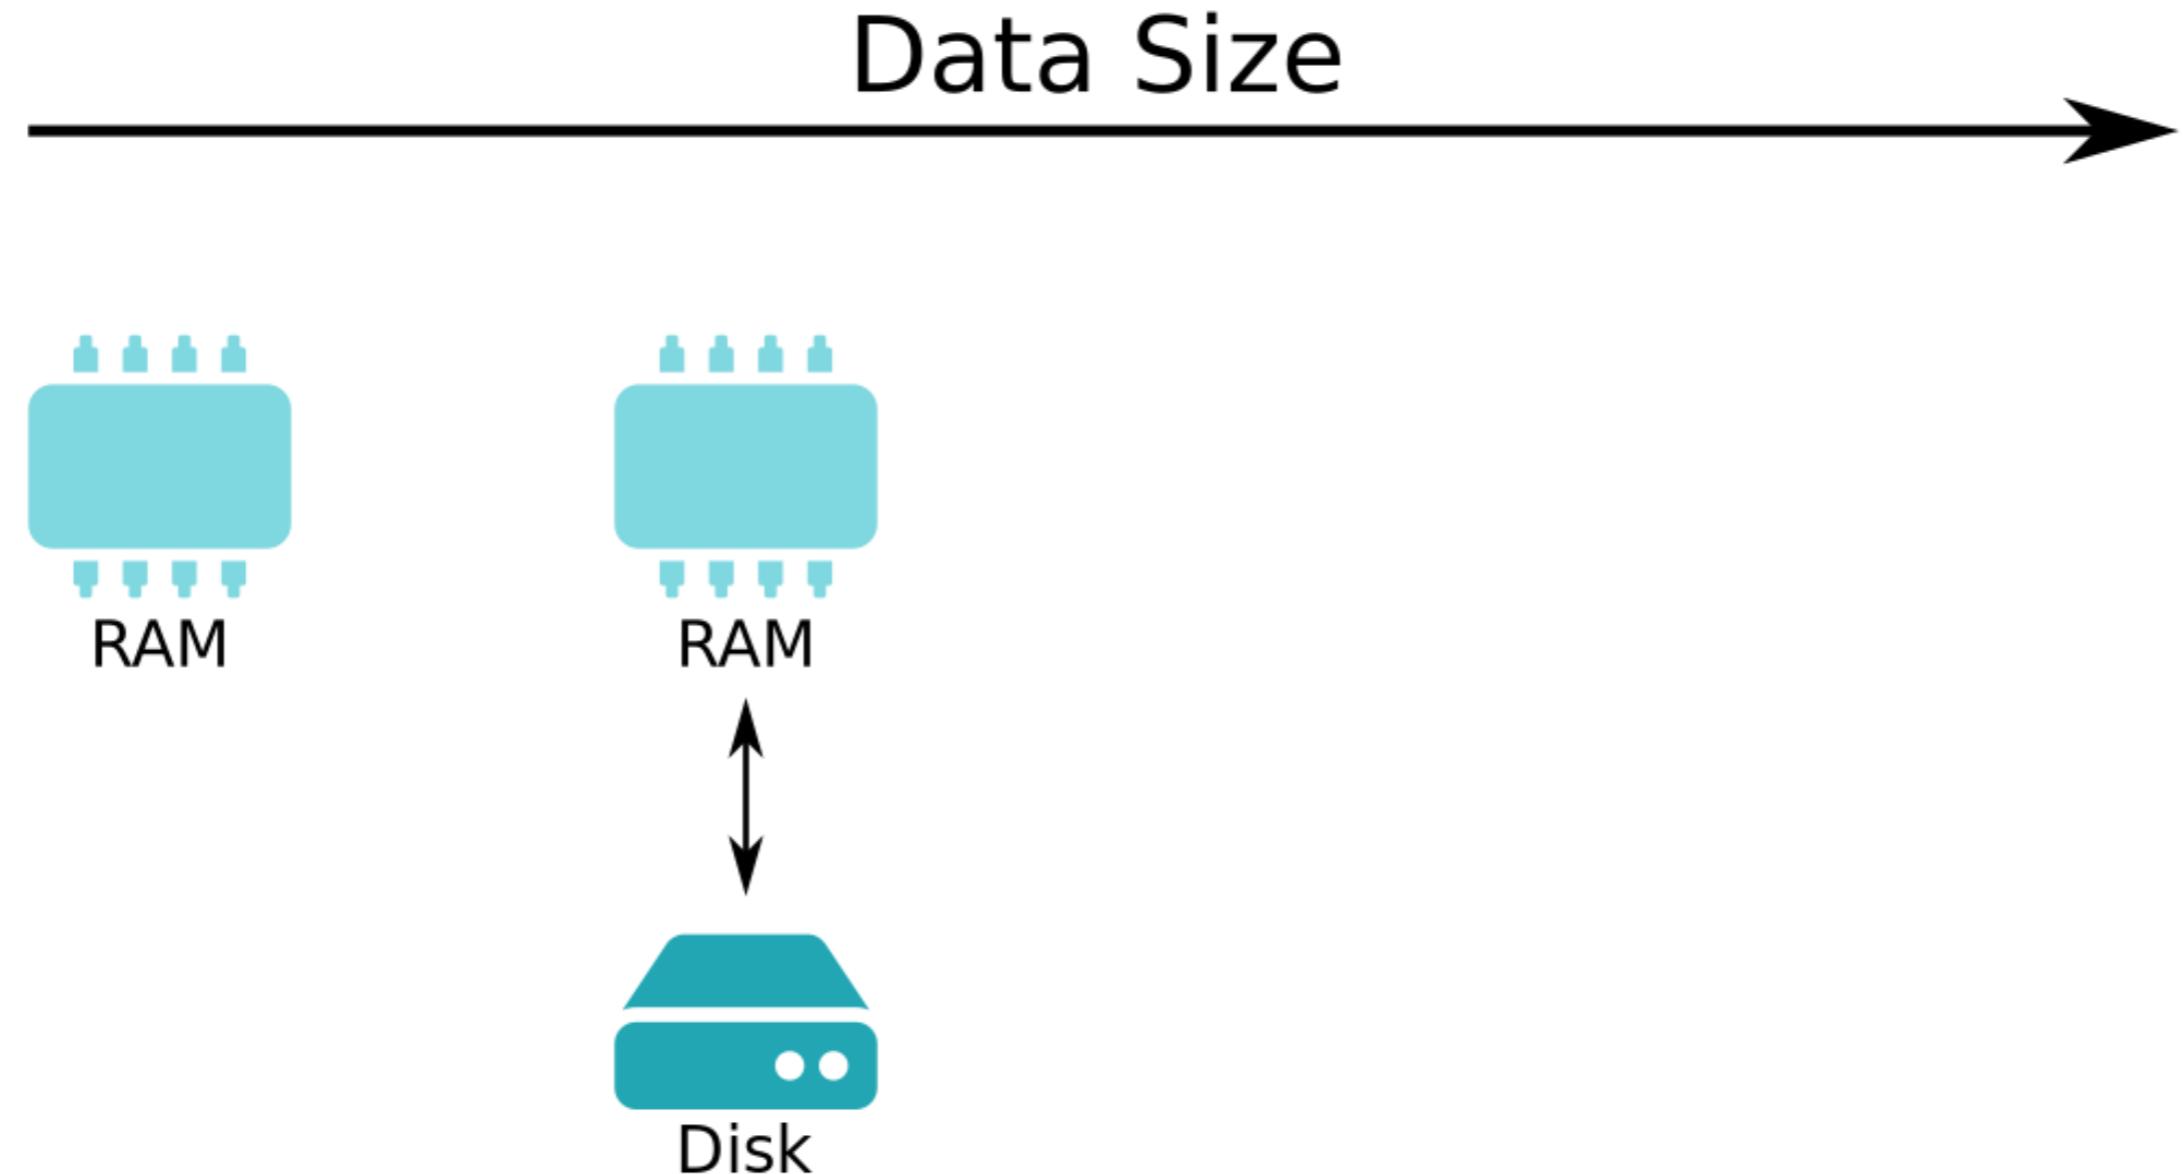
Classification



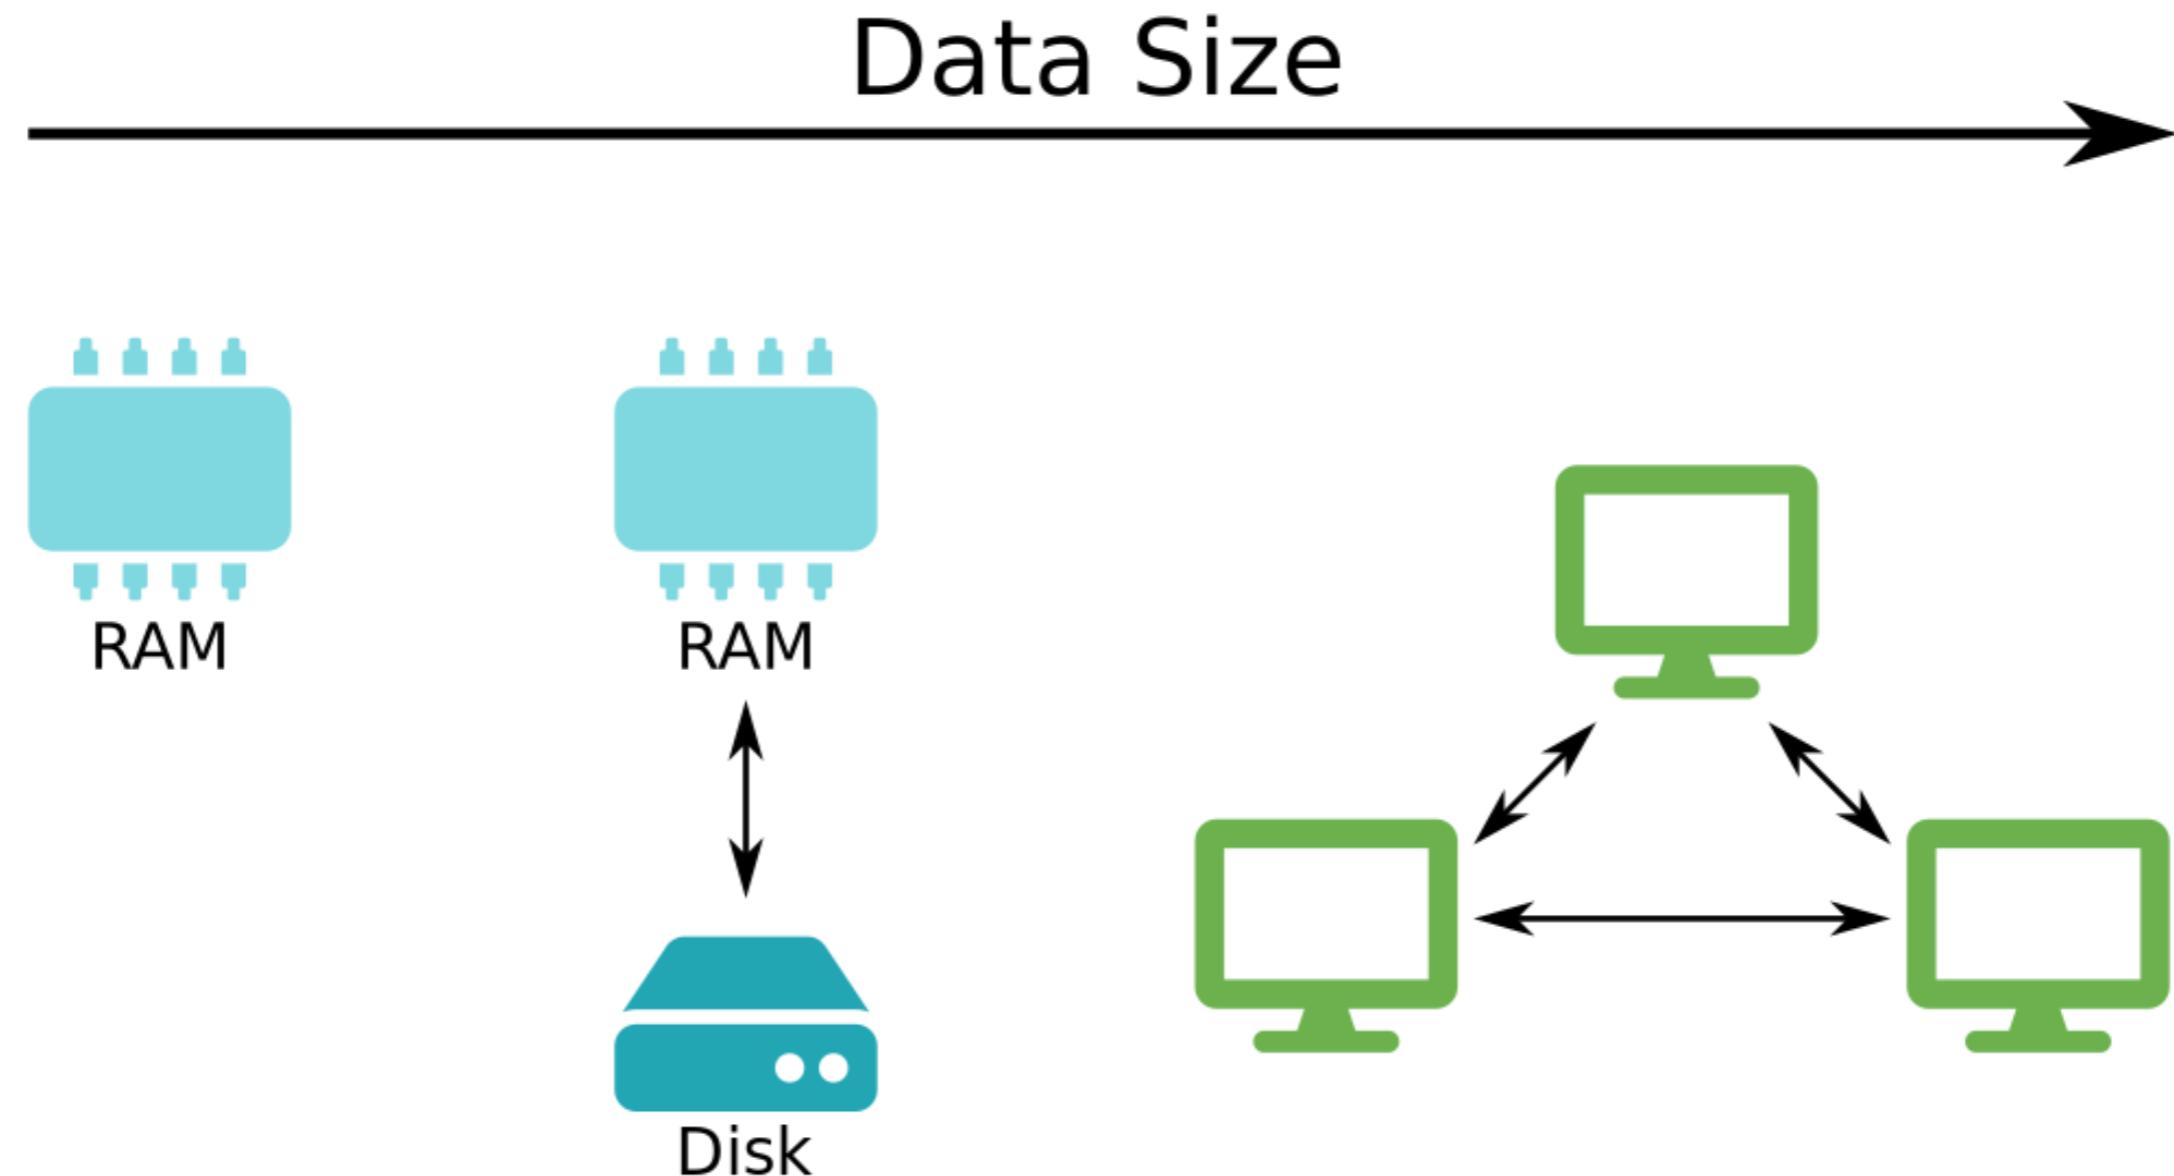
Data in RAM



Data exceeds RAM



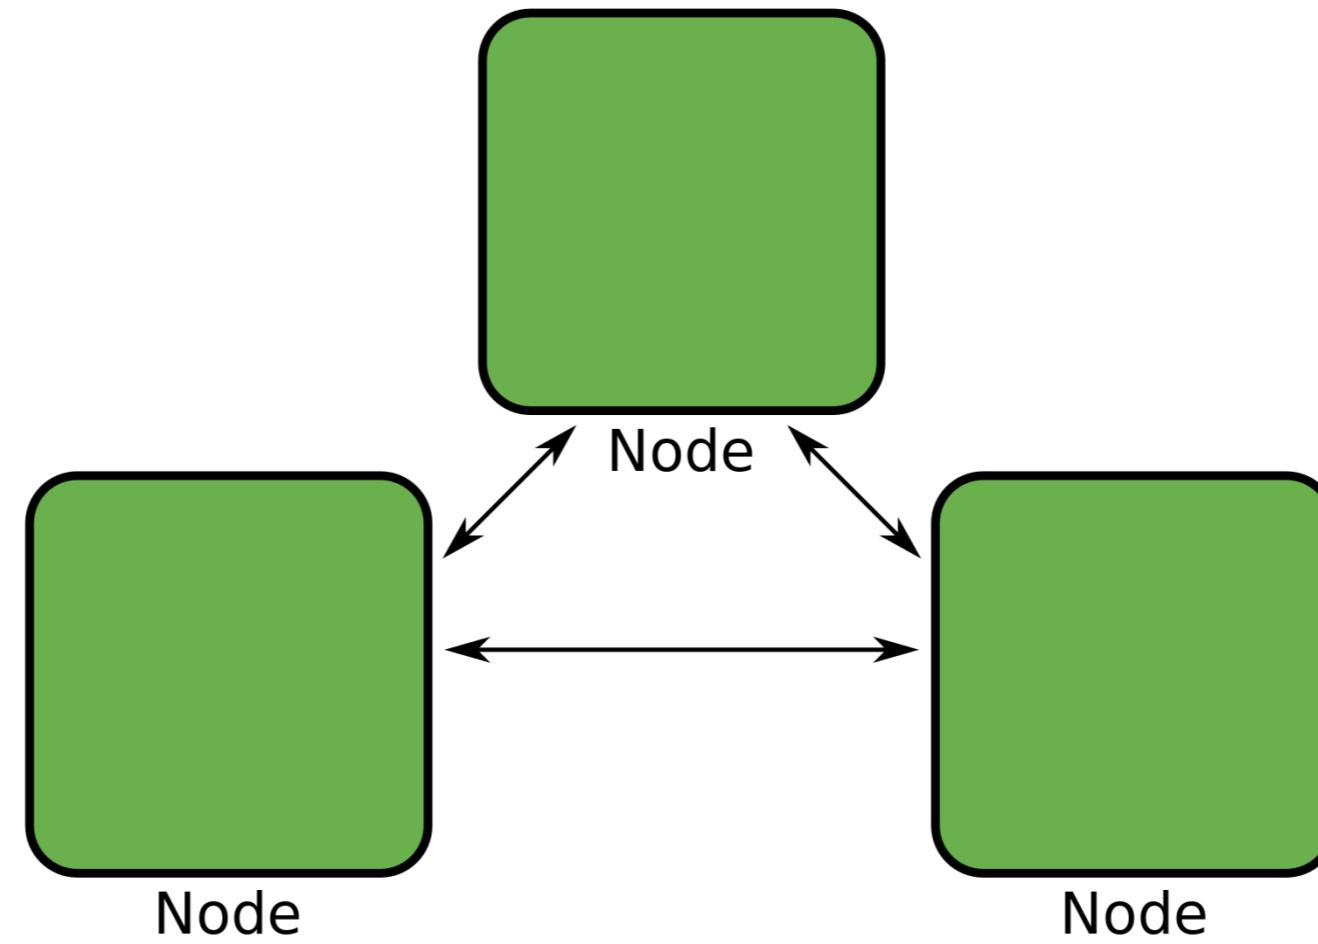
Data distributed across a cluster

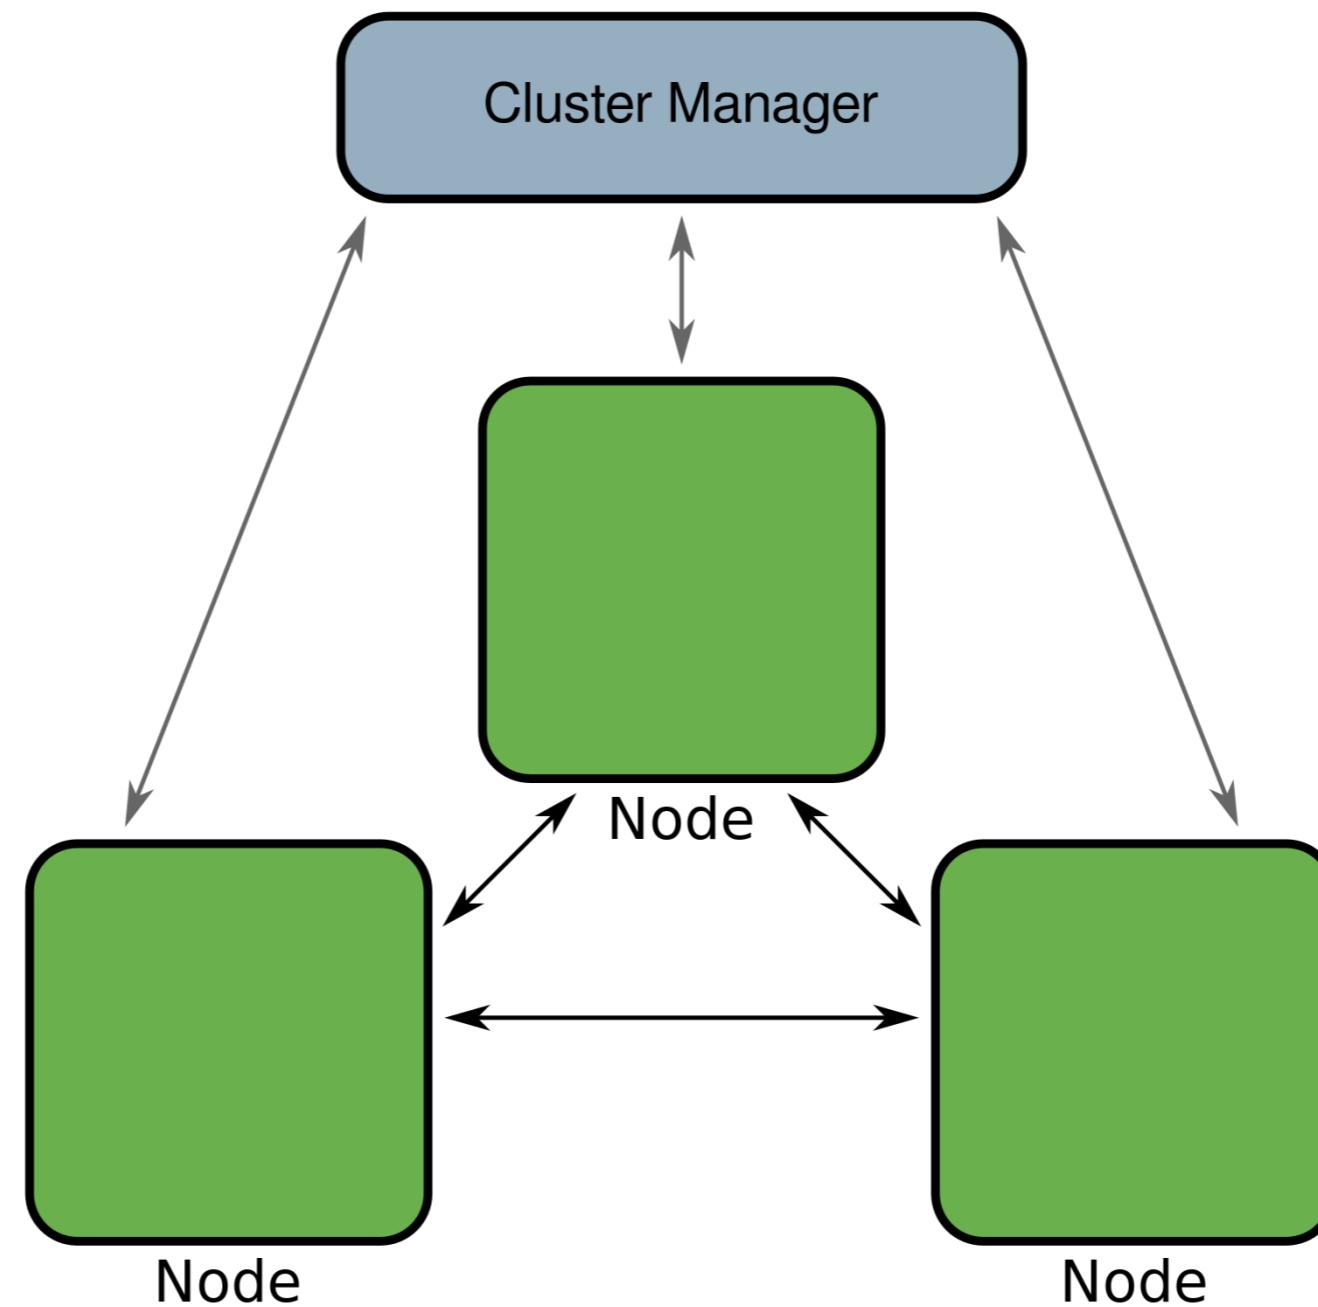


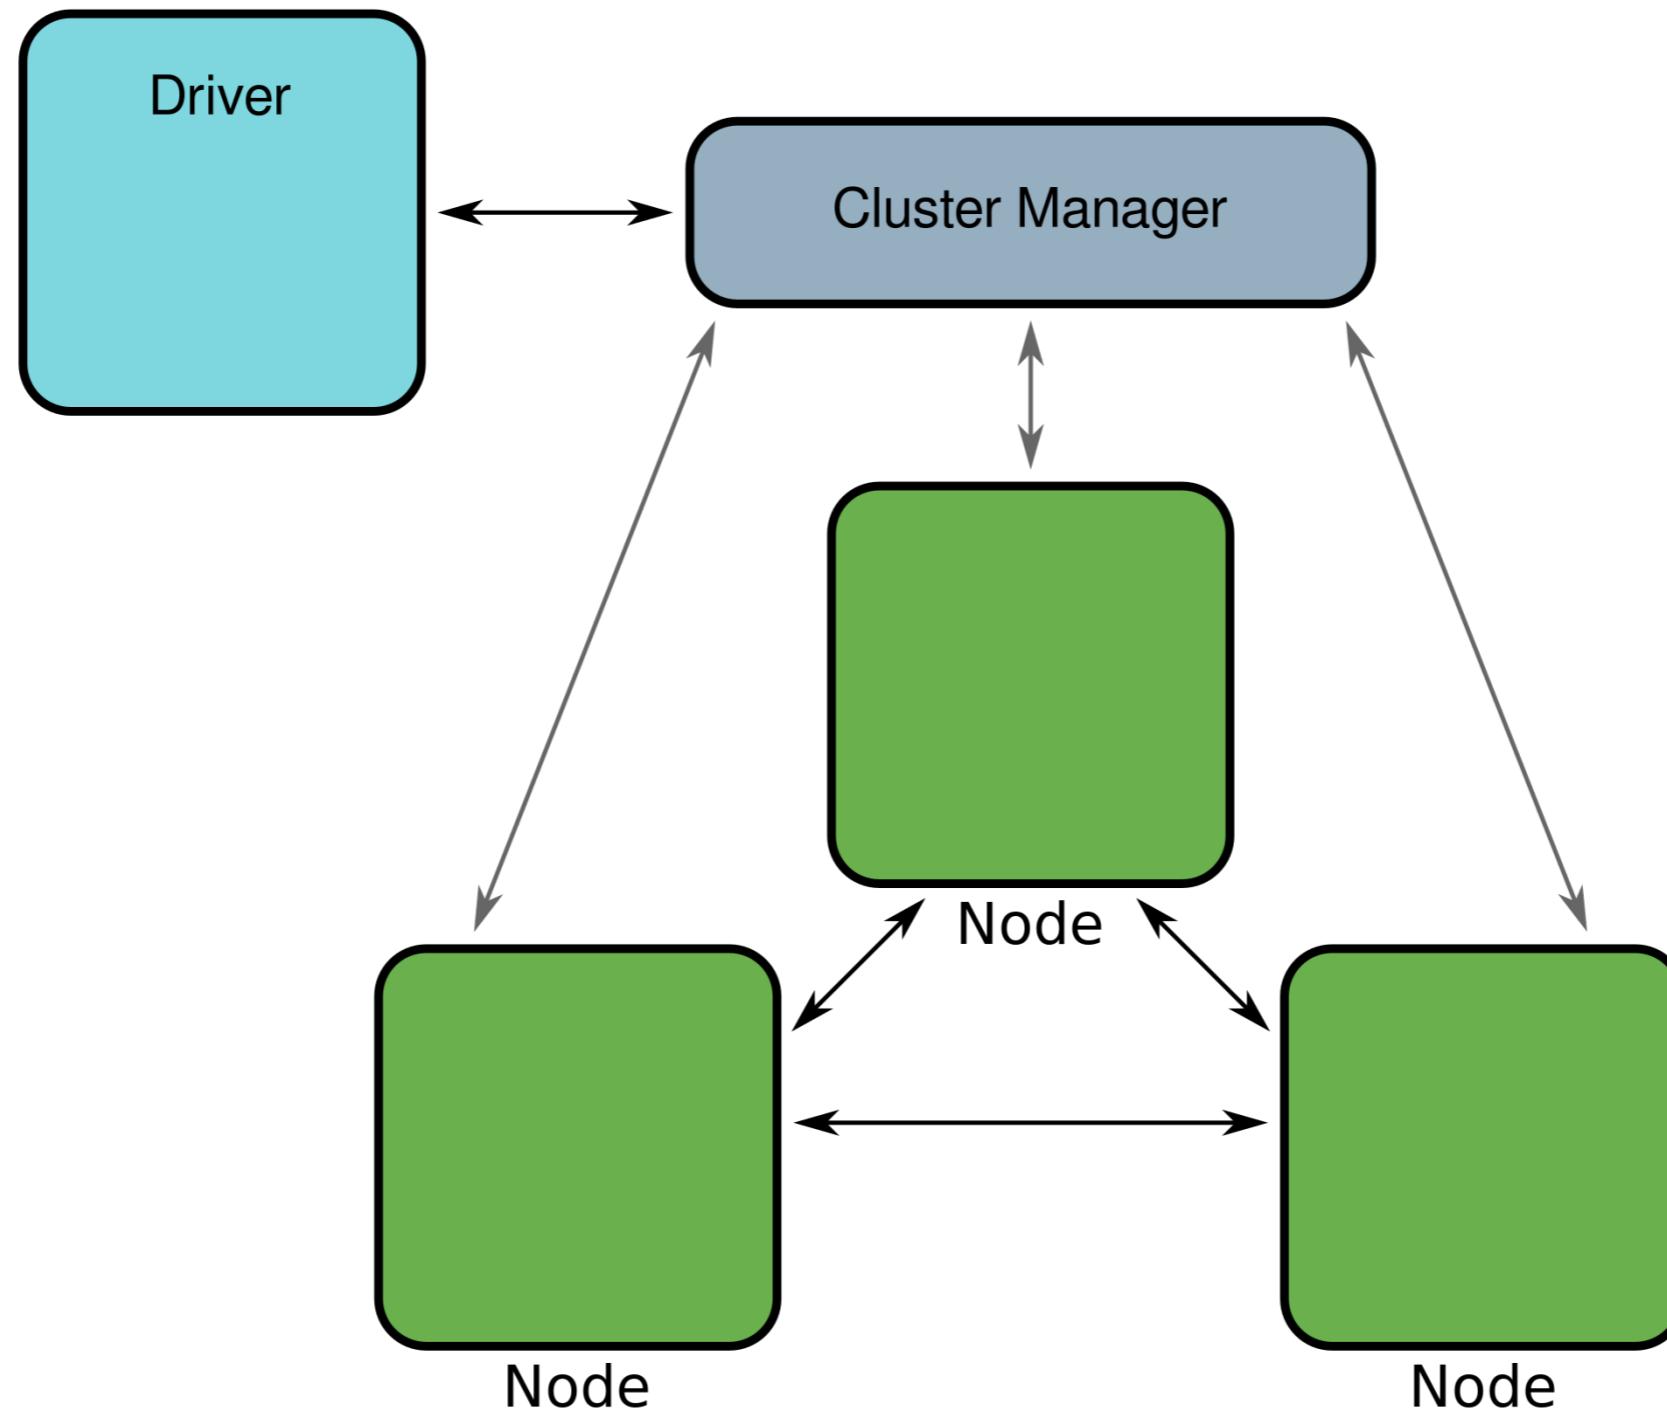
What is Spark?

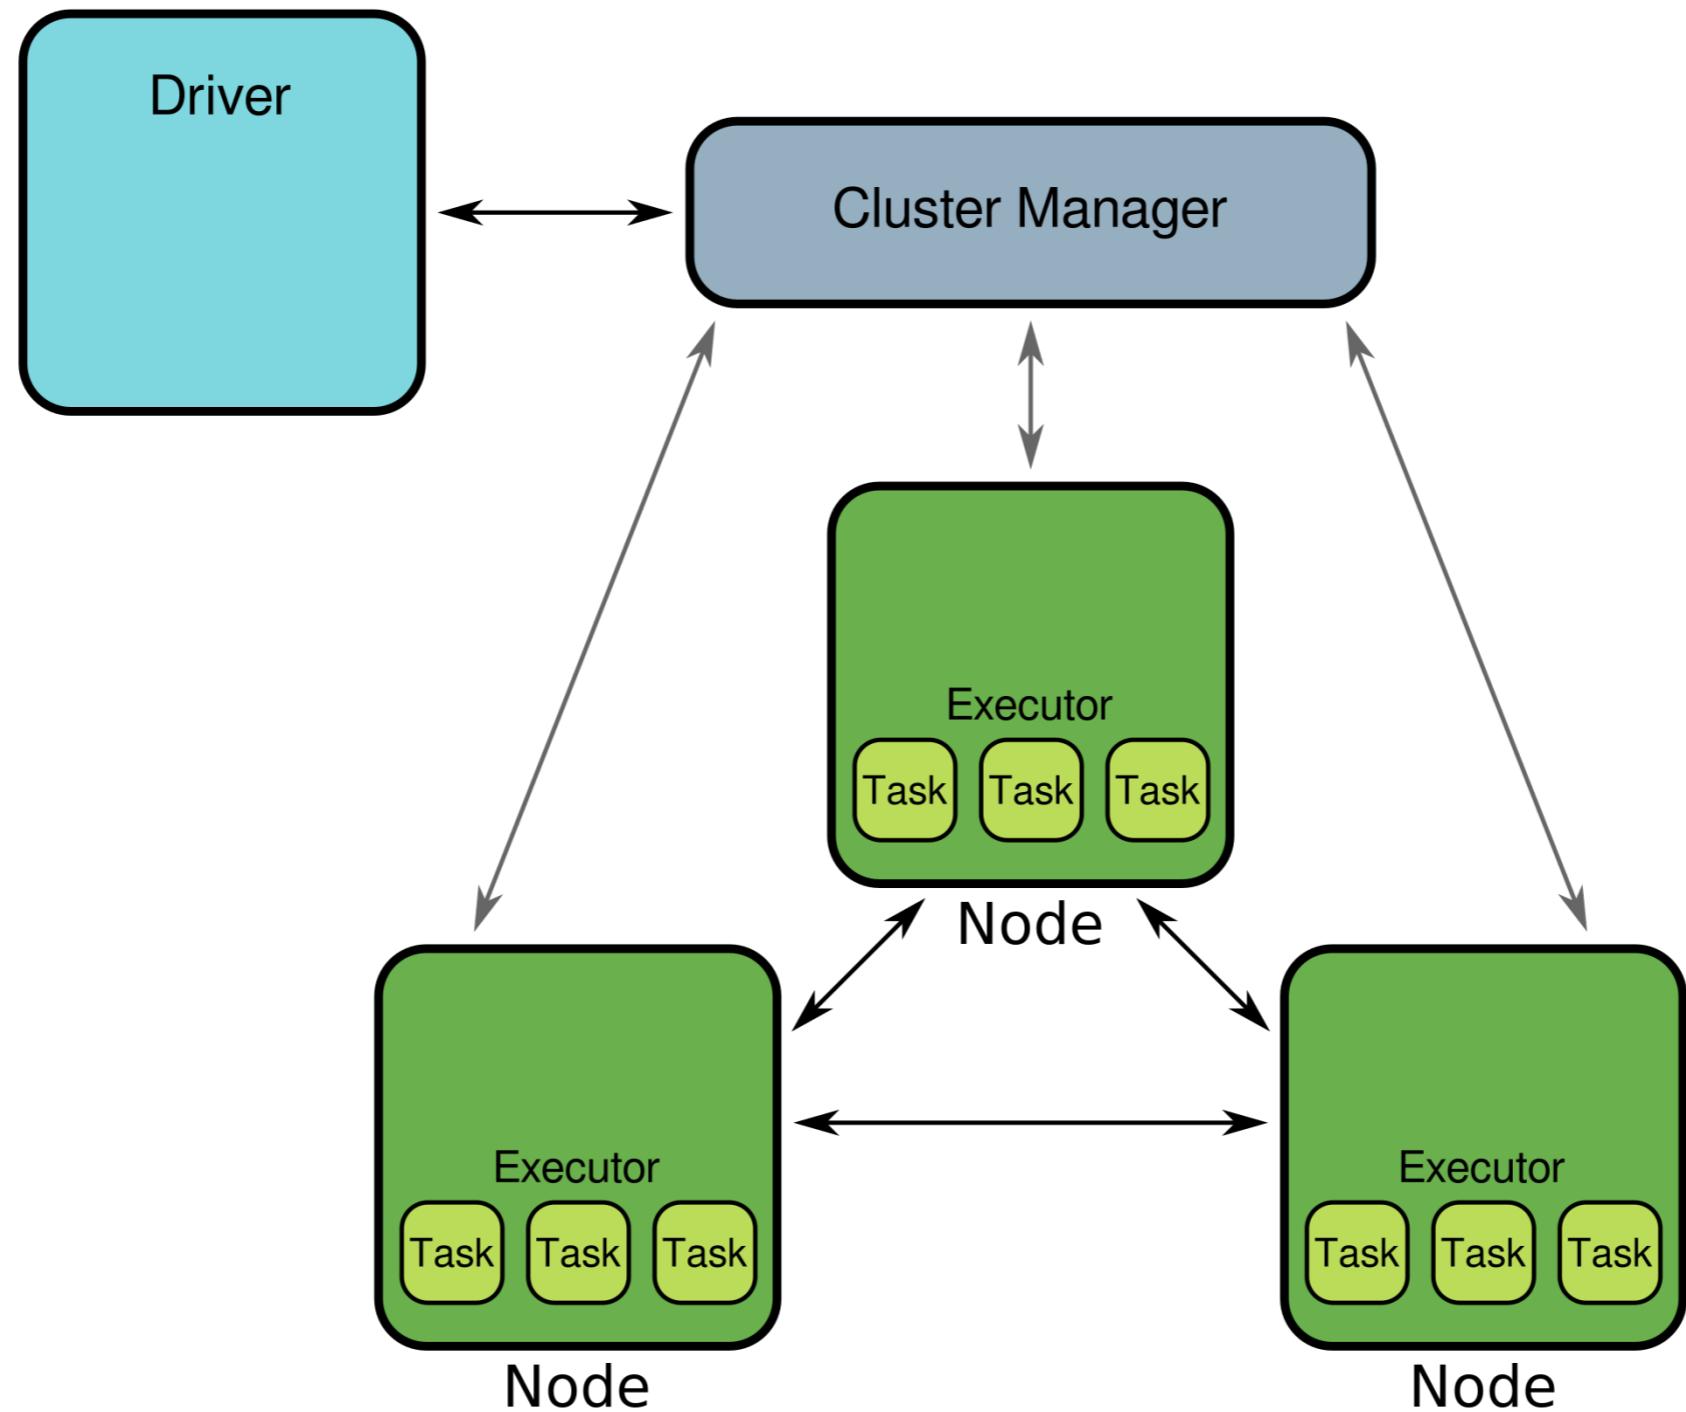


- Compute across a distributed **cluster**.
- Data processed in memory.
- Well documented high-level **API**.









Onward!

MACHINE LEARNING WITH PYSPARK

Connecting to Spark

MACHINE LEARNING WITH PYSPARK



Andrew Collier

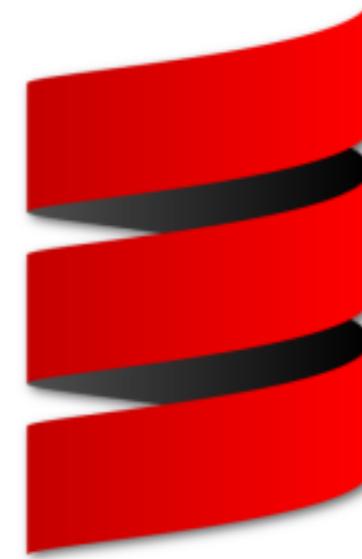
Data Scientist, Fathom Data

Interacting with Spark

Java



Scala



Python



R



Languages for interacting with Spark.

- Java — low-level, compiled
- Scala, Python and R — high-level with interactive REPL

Importing pyspark

From Python import the `pyspark` module.

```
import pyspark
```

Check version of the `pyspark` module.

```
pyspark.__version__
```

```
'2.4.1'
```

Sub-modules

In addition to `pyspark` there are

- Structured Data — `pyspark.sql`
- Streaming Data — `pyspark.streaming`
- Machine Learning — `pyspark.mllib` (deprecated) and `pyspark.ml`

Spark URL

Remote Cluster using Spark URL — `spark://<IP address | DNS name>:<port>`

Example:

- `spark://13.59.151.161:7077`

Local Cluster

Examples:

- `local` — only 1 core;
- `local[4]` — 4 cores; or
- `local[*]` — all available cores.

Creating a SparkSession

```
from pyspark.sql import SparkSession
```

Create a local cluster using a `SparkSession` builder.

```
spark = SparkSession.builder \  
    .master('local[*]') \  
    .appName('first_spark_application') \  
    .getOrCreate()
```

Interact with Spark...

```
# Close connection to Spark  
>>> spark.stop()
```

Let's connect to Spark!

MACHINE LEARNING WITH PYSPARK

Loading Data

MACHINE LEARNING WITH PYSPARK



Andrew Collier

Data Scientist, Fathom Data

DataFrames: A refresher

DataFrame for tabular data.

123	abc	123	abc
123	abc	123	abc
123	abc	123	abc
123	abc	123	abc
123	abc	123	abc

Selected methods:

- `count()`
- `show()`
- `printSchema()`

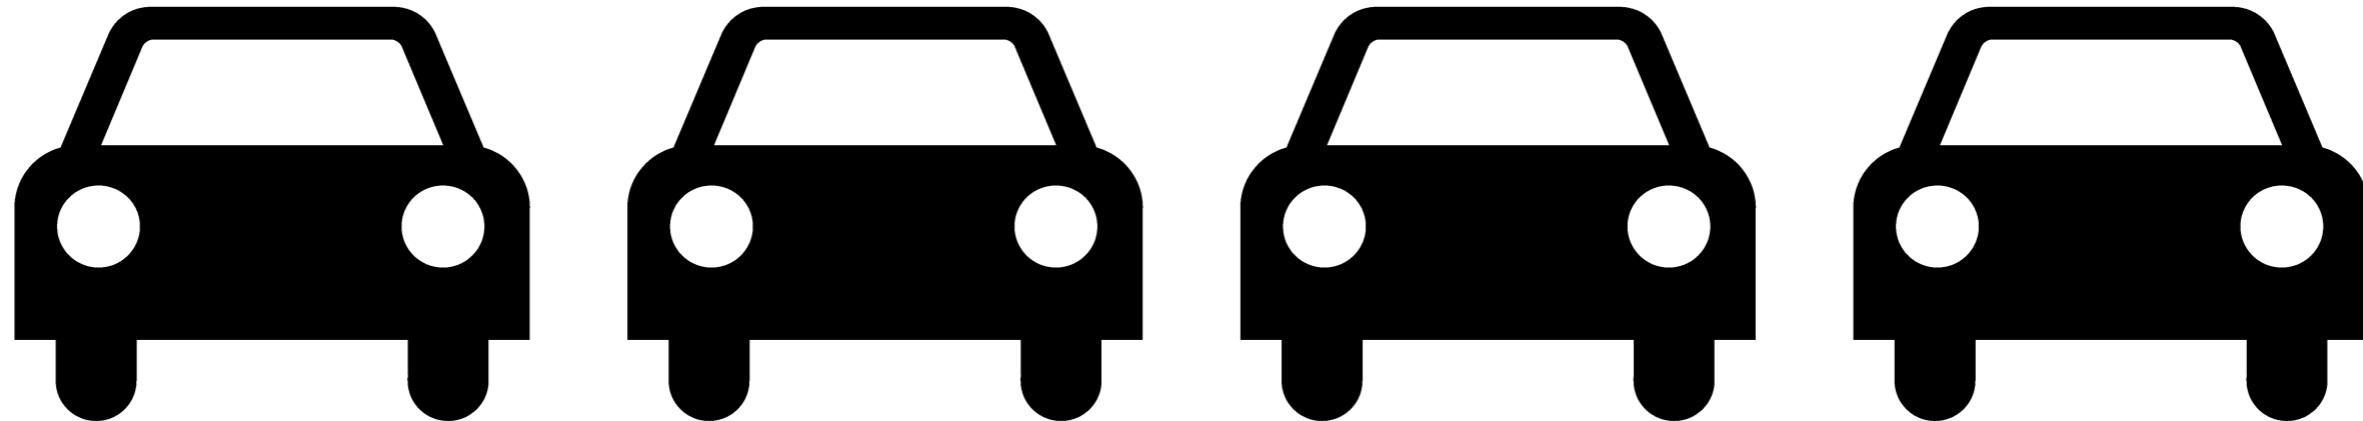
Selected attributes:

- `dtypes`

CSV data for cars

The first few lines from the 'cars.csv' file.

```
mfr,mod,org,type,cyl,size,weight,len,rpm,cons  
Mazda,RX-7,non-USA,Sporty,NA,1.3,2895,169,6500,9.41  
Nissan,Maxima,non-USA,Midsize,6,3,3200,188,5200,9.05  
Chevrolet,Cavalier,USA,Compact,4,2.2,2490,182,5200,6.53  
Subaru,Legacy,non-USA,Compact,4,2.2,3085,179,5600,7.84  
Ford,Escort,USA,Small,4,1.8,2530,171,6500,7.84
```



Reading data from CSV

The `.csv()` method reads a CSV file and returns a `DataFrame`.

```
cars = spark.read.csv('cars.csv', header=True)
```

Optional arguments:

- `header` — is first row a header? (default: `False`)
- `sep` — field separator (default: a comma `', '`)
- `schema` — explicit column data types
- `inferSchema` — deduce column data types from data?
- `nullValue` — placeholder for missing data

Peek at the data

The first five records from the `DataFrame`.

```
cars.show(5)
```

```
+-----+-----+-----+-----+-----+-----+-----+
|      mfr|      mod|      org|    type|cyl|size|weight|len| rpm|cons|
+-----+-----+-----+-----+-----+-----+-----+-----+
|  Mazda| RX-7|non-USA| Sporty| NA| 1.3| 2895|169|6500|9.41|
|  Nissan| Maxima|non-USA|Midsized| 6|   3| 3200|188|5200|9.05|
| Chevrolet|Cavalier| USA|Compact| 4| 2.2| 2490|182|5200|6.53|
|  Subaru| Legacy|non-USA|Compact| 4| 2.2| 3085|179|5600|7.84|
|   Ford| Escort| USA| Small| 4| 1.8| 2530|171|6500|7.84|
+-----+-----+-----+-----+-----+-----+-----+
```

Check column types

```
cars.printSchema()
```

```
root
|-- mfr: string (nullable = true)
|-- mod: string (nullable = true)
|-- org: string (nullable = true)
|-- type: string (nullable = true)
|-- cyl: string (nullable = true)
|-- size: string (nullable = true)
|-- weight: string (nullable = true)
|-- len: string (nullable = true)
|-- rpm: string (nullable = true)
|-- cons: string (nullable = true)
```

Inferring column types from data

```
cars = spark.read.csv("cars.csv", header=True, inferSchema=True)  
cars.dtypes
```

```
[('mfr', 'string'),  
 ('mod', 'string'),  
 ('org', 'string'),  
 ('type', 'string'),  
 ('cyl', 'string'),  
 ('size', 'double'),  
 ('weight', 'int'),  
 ('len', 'int'),  
 ('rpm', 'int'),  
 ('cons', 'double')]
```

Dealing with missing data

Handle missing data using the `nullValue` argument.

```
cars = spark.read.csv("cars.csv", header=True, inferSchema=True, nullValue='NA')
```

The `nullValue` argument is case sensitive.

Specify column types

```
schema = StructType([
    StructField("maker", StringType()),
    StructField("model", StringType()),
    StructField("origin", StringType()),
    StructField("type", StringType()),
    StructField("cyl", IntegerType()),
    StructField("size", DoubleType()),
    StructField("weight", IntegerType()),
    StructField("length", DoubleType()),
    StructField("rpm", IntegerType()),
    StructField("consumption", DoubleType())
])
cars = spark.read.csv("cars.csv", header=True, schema=schema, nullValue='NA')
```

Final cars data

maker	model	origin	type	cyl	size	weight	length	rpm	consumption
Mazda	RX-7	non-USA	Sporty	null	1.3	2895	169.0	6500	9.41
Nissan	Maxima	non-USA	Midsize	6	3.0	3200	188.0	5200	9.05
Chevrolet	Cavalier	USA	Compact	4	2.2	2490	182.0	5200	6.53
Subaru	Legacy	non-USA	Compact	4	2.2	3085	179.0	5600	7.84
Ford	Escort	USA	Small	4	1.8	2530	171.0	6500	7.84
Mercury	Capri	USA	Sporty	4	1.6	2450	166.0	5750	9.05
Oldsmobile	Cutlass Ciera	USA	Midsize	4	2.2	2890	190.0	5200	7.59
Saab	900	non-USA	Compact	4	2.1	2775	184.0	6000	9.05
Dodge	Caravan	USA	Van	6	3.0	3705	175.0	5000	11.2

Let's load some data!

MACHINE LEARNING WITH PYSPARK