

# One-Hot Encoding

## MACHINE LEARNING WITH PYSPARK



**Andrew Collier**

Data Scientist, Fathom Data

# The problem with indexed values

```
# Counts for 'type' category
```

```
+-----+-----+
|   type|count|
+-----+-----+
|Midsize|    22|
|  Small|    21|
|Compact|    16|
| Sporty|    14|
|  Large|    11|
|     Van|     9|
+-----+-----+
```

```
# Numerical indices for 'type' category
```

```
+-----+-----+
|   type|type_idx|
+-----+-----+
|Midsize|      0.0|
|  Small|      1.0|
|Compact|      2.0|
| Sporty|      3.0|
|  Large|      4.0|
|     Van|      5.0|
+-----+-----+
```

# Dummy variables

type	Midsized	Small	Compact	Sporty	Large	Van
Midsized	X					
Small		X				
Compact	==>		X			
Sporty				X		
Large					X	
Van						X

Each categorical level becomes a column.

# Dummy variables: binary encoding

type	Midsize	Small	Compact	Sporty	Large	Van
Midsize	1	0	0	0	0	0
Small	0	1	0	0	0	0
Compact	0	0	1	0	0	0
Sporty	0	0	0	1	0	0
Large	0	0	0	0	1	0
Van	0	0	0	0	0	1

Binary values indicate the presence (1) or absence (0) of the corresponding level.

# Dummy variables: sparse representation

type	Midsized	Small	Compact	Sporty	Large	Van	Column	Value
	1	0	0	0	0	0	0	1
Midsized	1	0	0	0	0	0	0	1
Small	0	1	0	0	0	0	0	1
Compact	0	0	1	0	0	0	0	1
Sporty	0	0	0	1	0	0	0	1
Large	0	0	0	0	1	0	0	1
Van	0	0	0	0	0	1	0	1

Sparse representation: store column index and value.

# Dummy variables: redundant column

type	Midsized	Small	Compact	Sporty	Large	Column	Value
	1	0	0	0	0	0	1
Midsized	1	0	0	0	0	0	1
Small	0	1	0	0	0	0	1
Compact	0	0	1	0	0	0	2
Sporty	0	0	0	1	0	0	3
Large	0	0	0	0	1	0	4
Van	0	0	0	0	0	0	0

Levels are mutually exclusive, so drop one.

# One-hot encoding

```
from pyspark.ml.feature import OneHotEncoder  
  
onehot = OneHotEncoder(inputCols=['type_idx'], outputCols=['type_dummy'])
```

Fit the encoder to the data.

```
onehot = onehot.fit(cars)
```

```
# How many category levels?  
onehot.categorySizes
```

[6]

# One-hot encoding

```
cars = onehot.transform(cars)
cars.select('type', 'type_idx', 'type_dummy').distinct().sort('type_idx').show()
```

```
+-----+-----+
| type|type_idx|  type_dummy|
+-----+-----+-----+
|Midsize|    0.0|(5,[0],[1.0])|
| Small|    1.0|(5,[1],[1.0])|
| Compact|   2.0|(5,[2],[1.0])|
| Sporty|   3.0|(5,[3],[1.0])|
| Large|   4.0|(5,[4],[1.0])|
| Van|    5.0|  (5,[],[])|
+-----+-----+
```

# Dense versus sparse

```
from pyspark.mllib.linalg import DenseVector, SparseVector
```

Store this vector: [1, 0, 0, 0, 0, 7, 0, 0].

```
DenseVector([1, 0, 0, 0, 0, 7, 0, 0])
```

```
DenseVector([1.0, 0.0, 0.0, 0.0, 0.0, 7.0, 0.0, 0.0])
```

```
SparseVector(8, [0, 5], [1, 7])
```

```
SparseVector(8, {0: 1.0, 5: 7.0})
```

# One-Hot Encode categoricals

MACHINE LEARNING WITH PYSPARK

# Regression


MACHINE LEARNING WITH PYSPARK




**Andrew Collier**

Data Scientist, Fathom Data


# Consumption versus mass: scatter




# Consumption versus mass: fit



# Consumption versus mass: alternative fits



# Consumption versus mass: residuals



# Loss function

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

MSE = "Mean Squared Error"

# Loss function: Observed values

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$y_i$  — observed values

# Loss function: Model values

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$y_i$  — observed values

$\hat{y}_i$  — model values

# Loss function: Mean

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$y_i$  — observed values

$\hat{y}_i$  — model values

# Assemble predictors

Predict `consumption` using `mass`, `cyl` and `type_dummy`.

Consolidate predictors into a single column.

mass	cyl	type_dummy	features	consumption
1451.0	6	(5,[0],[1.0])	(7,[0,1,2],[1451.0,6.0,1.0])	9.05
1129.0	4	(5,[2],[1.0])	(7,[0,1,4],[1129.0,4.0,1.0])	6.53
1399.0	4	(5,[2],[1.0])	(7,[0,1,4],[1399.0,4.0,1.0])	7.84
1147.0	4	(5,[1],[1.0])	(7,[0,1,3],[1147.0,4.0,1.0])	7.84
1111.0	4	(5,[3],[1.0])	(7,[0,1,5],[1111.0,4.0,1.0])	9.05

# Build regression model

```
from pyspark.ml.regression import LinearRegression  
  
regression = LinearRegression(labelCol='consumption')
```

Fit to `cars_train` (training data).


```
regression = regression.fit(cars_train)
```

Predict on `cars_test` (testing data).

```
predictions = regression.transform(cars_test)
```

# Examine predictions

consumption	prediction
7.84	8.92699470743403
9.41	9.379295891451353
8.11	7.23487264538364
9.05	9.409860194333735
7.84	7.059190923328711
7.84	7.785909738591766
7.59	8.129959405168547
5.11	6.836843743852942
8.11	7.17173702652015



# Calculate RMSE


```
from pyspark.ml.evaluation import RegressionEvaluator  
  
# Find RMSE (Root Mean Squared Error)  
RegressionEvaluator(labelCol='consumption').evaluate(predictions)
```

```
0.708699086182001
```

A `RegressionEvaluator` can also calculate the following metrics:

- `mae` (Mean Absolute Error)
- `r2` ( $R^2$ )
- `mse` (Mean Squared Error).

# Consumption versus mass: intercept



# Examine intercept


```
regression.intercept
```

```
4.9450616833727095
```

This is the fuel consumption in the (hypothetical) case that:

- `mass` = 0
- `cyl` = 0 and
- vehicle type is 'Van'.

# Consumption versus mass: slope



# Examine Coefficients

```
regression.coefficients
```

```
DenseVector([0.0027, 0.1897, -1.309, -1.7933, -1.3594, -1.2917, -1.9693])
```

```
mass      0.0027  
cyl       0.1897
```

```
Midsize   -1.3090  
Small     -1.7933  
Compact   -1.3594  
Sporty    -1.2917  
Large     -1.9693
```

# Regression for numeric predictions

MACHINE LEARNING WITH PYSPARK

# Bucketing & Engineering


MACHINE LEARNING WITH PYSPARK





**Andrew Collier**

Data Scientist, Fathom Data



# Bucketing





# Bucketing heights




# Bucketing heights



# Bucketing heights



# Bucketing heights




height	height_bin
1.42	short
1.45	short
1.47	short
1.50	short
1.52	average
1.57	average
1.60	average
1.75	average
1.85	tall
1.88	tall

# RPM histogram

Car RPM has "natural" breaks:

- $\text{RPM} < 4500$  — low
- $\text{RPM} > 6000$  — high
- otherwise — medium.



# RPM buckets

```
from pyspark.ml.feature import Bucketizer  
  
bucketizer = Bucketizer(splits=[3500, 4500, 6000, 6500],  
                        inputCol="rpm",  
                        outputCol="rpm_bin")
```

Apply buckets to `rpm` column.

```
bucketed = bucketizer.transform(cars)
```

# RPM buckets

```
bucketed.select('rpm', 'rpm_bin').show(5)
```

```
+----+-----+
| rpm|rpm_bin|
+----+-----+
|3800|    0.0|
|4500|    1.0|
|5750|    1.0|
|5300|    1.0|
|6200|    2.0|
+----+-----+
```

```
bucketed.groupBy('rpm_bin').count().show()
```

```
+----+-----+
|rpm_bin|count|
+----+-----+
|    0.0|    8| <- low
|    1.0|   67| <- medium
|    2.0|   17| <- high
+----+-----+
```

# One-hot encoded RPM buckets

The RPM buckets are one-hot encoded to dummy variables.

```
+-----+-----+
| rpm_bin|    rpm_dummy|
+-----+-----+
| 0.0|(2,[0],[1.0])| <- low
| 1.0|(2,[1],[1.0])| <- medium
| 2.0|(2,[],[])| <- high
+-----+-----+
```

The 'high' RPM bucket is the reference level and doesn't get a dummy variable.

# Model with bucketed RPM

regression.coefficients

```
DenseVector([1.3814, 0.1433])
```

```
+-----+  
| rpm_bin|      rpm_dummy|  
+-----+-----+  
| 0.0|(2,[0],[1.0])| <- low  
| 1.0|(2,[1],[1.0])| <- medium  
| 2.0| (2,[],[])| <- high  
+-----+
```

regression.intercept

```
8.1835
```

Consumption for 'low' RPM:

```
8.1835 + 1.3814 = 9.5649
```

Consumption for 'medium' RPM:

```
8.1835 + 0.1433 = 8.3268
```

# More feature engineering



Operations on a single column:

- `log()`
- `sqrt()`
- `pow()`


Operations on two columns:

- `product`
- `ratio.`

# Mass & Height to BMI



# Mass & Height to BMI



height	mass	bmi	bmi = mass / height^2
1.52	77.1	33.2	
1.60	58.1	22.7	
1.57	122.0	49.4	
1.75	95.3	31.0	
1.80	99.8	30.7	
1.65	90.7	33.3	
1.60	70.3	27.5	
1.78	81.6	25.8	
1.65	77.1	28.3	
1.78	128.0	40.5	

# Engineering density

```
cars = cars.withColumn('density_line', cars.mass / cars.length)      # Linear density
cars = cars.withColumn('density_quad', cars.mass / cars.length**2)    # Area density
cars = cars.withColumn('density_cube', cars.mass / cars.length**3)    # Volume density
```

```
+-----+-----+-----+-----+
| mass|length|density_line|density_quad|density_cube|
+-----+-----+-----+-----+
|1451.0| 4.775|303.87434554|63.638606397|13.327456837|
|1129.0| 4.623|244.21371403|52.825808790|11.426737787|
|1399.0| 4.547|307.67539036|67.665579583|14.881367843|
+-----+-----+-----+-----+
```

# **Let's engineer some features!**

MACHINE LEARNING WITH PYSPARK

# Regularization


MACHINE LEARNING WITH PYSPARK





**Andrew Collier**

Data Scientist, Fathom Data


# Features: Only a few



# Features: Too many



# Features: Selected



# Loss function (revisited)

Linear regression aims to minimise the MSE.

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

# Loss function with regularization

Linear regression aims to minimise the MSE.

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 + \lambda f(\beta)$$

Add a *regularization* term which depends on coefficients.

# Regularization term

An extra *regularization* term is added to the loss function.

The regularization term can be either

- *Lasso* — absolute value of the coefficients
- *Ridge* — square of the coefficients

It's also possible to have a blend of Lasso and Ridge regression.

Strength of regularization determined by parameter  $\lambda$ :

- $\lambda = 0$  — no regularization (standard regression)
- $\lambda = \infty$  — complete regularization (all coefficients zero)

# Cars again

```
assembler = VectorAssembler(inputCols=[  
    'mass', 'cyl', 'type_dummy', 'density_line', 'density_quad', 'density_cube'  
], outputCol='features')  
cars = assembler.transform(cars)
```

```
+-----+-----+  
| features | consumption |  
+-----+-----+  
|[1451.0,6.0,1.0,0.0,0.0,0.0,303.8743455497,63.63860639785,13.32745683724]|9.05 |  
|[1129.0,4.0,0.0,0.0,1.0,0.0,0.0,244.2137140385,52.82580879050,11.42673778726]|6.53 |  
|[1399.0,4.0,0.0,0.0,1.0,0.0,0.0,307.6753903672,67.66557958374,14.88136784335]|7.84 |  
|[1147.0,4.0,0.0,1.0,0.0,0.0,0.0,264.1031545014,60.81122599620,14.00212433714]|7.84 |  
+-----+-----+
```

# Cars: Linear regression

Fit a (standard) Linear Regression model to the training data.

```
regression = LinearRegression(labelCol='consumption').fit(cars_train)
```

```
# RMSE on testing data  
0.708699086182001
```

Examine the coefficients:

```
regression.coefficients
```

```
DenseVector([-0.012, 0.174, -0.897, -1.445, -0.985, -1.071, -1.335, 0.189, -0.780, 1.160])
```

# Cars: Ridge regression

```
# alpha = 0 | lambda = 0.1 -> Ridge  
ridge = LinearRegression(labelCol='consumption', elasticNetParam=0, regParam=0.1)  
ridge.fit(cars_train)
```

```
# RMSE  
0.724535609745491
```

```
# Ridge coefficients  
DenseVector([ 0.001, 0.137, -0.395, -0.822, -0.450, -0.582, -0.806, 0.008, 0.029, 0.001])  
# Linear Regression coefficients  
DenseVector([-0.012, 0.174, -0.897, -1.445, -0.985, -1.071, -1.335, 0.189, -0.780, 1.160])
```

# Cars: Lasso regression

```
# alpha = 1 | lambda = 0.1 -> Lasso  
lasso = LinearRegression(labelCol='consumption', elasticNetParam=1, regParam=0.1)  
lasso.fit(cars_train)
```

```
# RMSE  
0.771988667026998
```

```
# Lasso coefficients  
DenseVector([ 0.0, 0.0, 0.0, -0.056, 0.0, 0.0, 0.0, 0.026, 0.0, 0.0])  
# Ridge coefficients  
DenseVector([ 0.001, 0.137, -0.395, -0.822, -0.450, -0.582, -0.806, 0.008, 0.029, 0.001])  
# Linear Regression coefficients  
DenseVector([-0.012, 0.174, -0.897, -1.445, -0.985, -1.071, -1.335, 0.189, -0.780, 1.160])
```

**Regularization →  
simple model**

MACHINE LEARNING WITH PYSPARK